

MACHINE LEARNING

Kernel regression and gaussian processes

Corso di Laurea Magistrale in Informatica

Università di Roma Tor Vergata

Prof. Giorgio Gambosi

a.a. 2021-2022




Kernel regression

- ⊙ In kernel regression methods, the target value corresponding to any item \mathbf{x} is predicted by referring to items in the training set, and in particular to the items which are closer to \mathbf{x} .
- ⊙ This is controlled by referring to a **kernel** function $\kappa_h(\mathbf{x})$, which is non zero only in an interval around 0
- ⊙ h is the **bandwidth** of the kernel, which controls the width of $\kappa_h(\mathbf{x})$

A possible, common kernel, is the gaussian (or RBF) kernel

$$g(\mathbf{x}) = e^{-\frac{|\mathbf{x}|^2}{2h^2}}$$



In regression, we are interested in estimating the conditional expectation

$$f(\mathbf{x}) = E[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt = \int t \frac{p(\mathbf{x}, t)}{p(\mathbf{x})} dt = \frac{\int t p(\mathbf{x}, t) dt}{p(\mathbf{x})} = \frac{\int t p(\mathbf{x}, t) dt}{\int p(\mathbf{x}, t) dt}$$

The joint distribution $p(\mathbf{x}, t)$ is approximated by means of a kernel function as

$$p(\mathbf{x}, t) \approx \frac{1}{n} \sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i) \kappa_h(t - t_i)$$

This results into

$$f(\mathbf{x}) = \frac{\int t \frac{1}{n} \sum_{i=1}^n \kappa_t(\mathbf{x} - \mathbf{x}_i) \kappa_h(t - t_i) dt}{\int \frac{1}{n} \sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i) \kappa_h(t - t_i) dt} = \frac{\sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i) \int t \kappa_h(t - t_i) dy}{\sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i) \int \kappa_h(t - t_i) dt}$$

and, since $\int \kappa_h(t - t_i) dt = 1$ and $\int t \kappa_h(t - t_i) dt = t_i$, we get

$$f(\mathbf{x}) = \frac{\sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i) t_i}{\sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i)}$$

By setting

$$w_i(\mathbf{x}) = \frac{\kappa_h(\mathbf{x} - \mathbf{x}_i)}{\sum_{j=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_j)}$$

we can write

$$f(\mathbf{x}) = \sum_{i=1}^n w_i(\mathbf{x}) t_i$$

that is, the predicted value is computed as a normalized linear combination of all target values, weighted by kernels (Nadaraya-Watson)

Locally weighted regression

In Nadaraya-Watson model, the prediction is performed by means of a normalized weighted combination of constant values (target values in the training set).

Locally weighted regression (LOESS) improves that approach by referring to a weighted version of the sum of squared differences loss function used in regression.

If a value t has to be predicted for an item \mathbf{x} , a “local” version of the loss function is considered, with weight $\kappa_i(\mathbf{x})$.

$$L(\mathbf{x}) = \sum_{i=1}^n \kappa_i(\mathbf{x})(\mathbf{w}^T \bar{\mathbf{x}}_i - t_i)^2 = \sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i)(\mathbf{w}^T \bar{\mathbf{x}}_i - t_i)^2$$

Weights $\kappa_i(\mathbf{x})$ are dependent from the “distance” between \mathbf{x} and \mathbf{x}_i , as measured by the kernel function

$$\kappa_i(\mathbf{x}) = \kappa_h(\mathbf{x} - \mathbf{x}_i)$$

The minimization of this loss function

$$\hat{\mathbf{w}}(\mathbf{x}) = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n \kappa_i(\mathbf{x}) (\mathbf{w}^T \bar{\mathbf{x}}_i - t_i)^2$$

has solution

$$\hat{\mathbf{w}}(\mathbf{x}) = (\bar{\mathbf{X}}^T \Psi(\mathbf{x}) \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \Psi(\mathbf{x}) \mathbf{t}$$

where $\Psi(\mathbf{x})$ is a diagonal $n \times n$ matrix with $\Psi(\mathbf{x})_{ii} = \kappa_i(\mathbf{x})$.

The prediction is then performed as usual, as

$$y = \hat{\mathbf{w}}(\mathbf{x})^T \bar{\mathbf{x}}$$

The same approach applied in the case of local regression can be applied for classification, by defining a weighted loss function to be minimized, with weights dependent from the item whose target must be predicted.

In this case, a weighted version of the cross entropy function is considered, which has to be maximized

$$L(\mathbf{x}) = \sum_{i=1}^n \kappa_h(\mathbf{x} - \mathbf{x}_i) (t_i \log p_i - (1 - t_i) \log(1 - p_i))$$

with $p_i = \sigma(\mathbf{w}^T \bar{\mathbf{x}}_i)$, as usual.

The loss function minimization can be performed, for example, by applying a suitable modification of the IRLS algorithm for logistic regression

Recap: some properties of Gaussian distribution

In order to introduce Gaussian processes and how they can be exploited for regression, let us first provide a short reminder on some properties of multivariate gaussian distributions.

Let $\mathbf{x} = (x_1, \dots, x_n)^T$ be a random vector with gaussian distribution $p(\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}, \Sigma)$ and let $\mathbf{x} = (\mathbf{x}_A, \mathbf{x}_B)$ be a partition of the components \mathbf{x} such that:

$$\odot \mathbf{x}_A = (x_1, \dots, x_r)^T$$

$$\odot \mathbf{x}_B = (x_{r+1}, \dots, x_n)^T$$

Then, the **marginal** densities $p(\mathbf{x}_A)$ and $p(\mathbf{x}_B)$ are both gaussian with means $\boldsymbol{\mu}_A, \boldsymbol{\mu}_B$ and covariance matrices Σ_A, Σ_B which can be derived from $\boldsymbol{\mu}, \Sigma$ by observing that

$$\boldsymbol{\mu} = (\boldsymbol{\mu}_A, \boldsymbol{\mu}_B)^T$$

$$\Sigma = \begin{pmatrix} \Sigma_A & \Sigma_{AB} \\ \Sigma_{AB}^T & \Sigma_B \end{pmatrix}$$

Recap: some properties of the Gaussian distribution

In the same situation, the conditional densities $p(\mathbf{x}_A|\mathbf{x}_B)$ and $p(\mathbf{x}_B|\mathbf{x}_A)$ are also gaussian with means

$$\boldsymbol{\mu}_{A|B} = \boldsymbol{\mu}_A + \Sigma_{AB}\Sigma_B^{-1}(\mathbf{x}_B - \boldsymbol{\mu}_B)$$

$$\boldsymbol{\mu}_{B|A} = \boldsymbol{\mu}_B + \Sigma_{BA}\Sigma_A^{-1}(\mathbf{x}_A - \boldsymbol{\mu}_A)$$

and covariance matrices

$$\Sigma_{A|B} = \Sigma_A - \Sigma_{AB}\Sigma_B^{-1}\Sigma_{BA}$$

$$\Sigma_{B|A} = \Sigma_B - \Sigma_{BA}\Sigma_A^{-1}\Sigma_{AB}$$

- ⊙ Multivariate gaussians on random vectors are useful for modeling finite collections of real-valued variables. They have nice analytical properties (see previous slides).
- ⊙ Gaussian processes: extension of multivariate gaussians to infinite-sized collections of real-valued variables.
- ⊙ We may think of gaussian processes as distributions not just over random vectors but over random real functions.

Let us first consider the case of functions defined over finite vectors.


- ⊙ Let $\chi = (\mathbf{x}_1, \dots, \mathbf{x}_m)$ be any finite vector, and let \mathcal{H} be the set of functions $f : \chi \mapsto \mathbb{R}$: f assigns a value $f(\mathbf{x}_i)$ to each $\mathbf{x}_i \in \chi$
 - A function $f \in \mathcal{H}$ can be described by the vector $(f(\mathbf{x}_1), \dots, f(\mathbf{x}_m))$
 - Any vector (y_1, \dots, y_m) can be seen as the description of a function $f \in \mathcal{H}$ such that $f(\mathbf{x}_i) = y_i$
 - The set \mathcal{H} is then in 1-to-1 correspondence with the set of vectors in \mathbb{R}^m
- ⊙ A probability distribution $p(\mathbf{x}), \mathbf{x} \in \mathbb{R}^m$ over m -dimensional real vectors is also a distribution $p(f), f \in \mathcal{H}$ over functions from \mathbb{R}^m to \mathbb{R}

Assume that $p(\mathbf{x})$ (or, equivalently, $p(f)$) is a (multivariate, m -dimensional) Gaussian distribution centered on $\mathbf{0}$ and with diagonal covariance $\sigma^2 \mathbf{I}$, that is

$$p(f|\sigma^2) = \mathcal{N}(f|\mathbf{0}, \sigma^2 \mathbf{I}) = \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{f(\mathbf{x}_i)^2}{2\sigma^2}}$$

- ⊙ This is equivalent to assuming that each function value $f(\mathbf{x}_i)$ has normal distribution with mean $\mathbf{0}$ and variance σ^2 , and that items are independent
- ⊙ A dependence between function values at different points could be modeled through a non-diagonal covariance matrix

Gaussian distributions over functions with finite domains



We may consider $p(f|\sigma^2)$ as a **prior** distribution of functions, with respect to the observation of the value t_j actually taken by any variable \mathbf{x}_j , $1 \leq j \leq m$.

Gaussian distributions over functions with finite domains

- ⊙ Assume now that for some subset $\mathbf{X} = \{\mathbf{x}_{i_1}, \dots, \mathbf{x}_{i_k}\}$ of component indices, the corresponding targets $\mathbf{t} = \{t_{i_1}, \dots, t_{i_k}\}$ are available.
- ⊙ the posterior distribution $p(f|\mathbf{X}, \mathbf{t})$ of functions (wrt to \mathbf{X}, \mathbf{t}) can be defined and derived according to Bayes' rule, provided a likelihood model is defined

$$p(\mathbf{X}, \mathbf{t}|f) = \prod_{\mathbf{x}_i \in \mathbf{X}} p(\mathbf{x}_i, t_i|f) = \prod_{\mathbf{x}_i \in \mathbf{X}} p(t_i|\mathbf{x}_i, f)p(\mathbf{x}_i|f) \propto \prod_{\mathbf{x}_i \in \mathbf{X}} p(t_i|\mathbf{x}_i, f)$$

- ⊙ for example, we could assume the usual likelihood $p(t|\mathbf{x}, f, \beta) = \mathcal{N}(t|f(\mathbf{x}), \beta)$, which implies

$$p(\mathbf{X}, \mathbf{t}|f, \beta) \propto \prod_{\mathbf{x}_i \in \mathbf{X}} \mathcal{N}(t_i|f(\mathbf{x}_i), \beta)$$

- ⊙ the posterior distribution then would be

$$p(f|\mathbf{X}, \mathbf{t}, \beta, \sigma^2) \propto \prod_{\mathbf{x}_i \in \mathbf{X}} \mathcal{N}(t_i|f(\mathbf{x}_i), \beta)p(f|\sigma^2)$$

Since both the prior and the posterior distributions of f are gaussian, the predictive distribution

$$p(t|\mathbf{x}, \mathbf{X}, \mathbf{t}, \beta, \sigma^2) = \int p(t|\mathbf{x}, f, \beta) p(f|\mathbf{X}, \mathbf{t}, \beta, \sigma^2) df$$

is itself a gaussian.

That would be the case also in the more general case when some dependency between function points is assumed. In this case, a general covariance matrix is defined for the prior distribution

$$p(f|\Sigma) = \mathcal{N}(f|\mathbf{0}, \Sigma)$$

- ⊙ In the case of infinite χ , we have to deal with an infinite collection of random variables.
- ⊙ In this case, the role of multidimensional distributions is covered by stochastic processes.
 - A *stochastic process* is a collection of random variables, $\{f(\mathbf{x}) : \mathbf{x} \in \chi\}$, indexed by elements from some set χ , known as the index set.
- ⊙ A **Gaussian process** is a stochastic process such that **for any** finite subset $\mathbf{x}_1, \dots, \mathbf{x}_n$ of χ , the function values $f(\mathbf{x}_1), \dots, f(\mathbf{x}_n)$ have joint multivariate Gaussian distribution

⊙ In order to define a Gaussian process, both a mean and a covariance function must be defined.

- a mean function $m : \mathbb{R}^d \mapsto \mathbb{R}$ mapping each point $\mathbf{x}_i \in \chi$ to the expectation

$$m(\mathbf{x}_i) = E_f[f(\mathbf{x}_i)]$$

of $f(\mathbf{x})$ over all functions f

- a covariance function $\kappa : \mathbb{R}^d \times \mathbb{R}^d \mapsto \mathbb{R}$ mapping each pair of variables $(\mathbf{x}_i, \mathbf{x}_j) \in \chi^2$ to the covariance

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = E_f[(f(\mathbf{x}_i) - m(\mathbf{x}_i))^2(f(\mathbf{x}_j) - m(\mathbf{x}_j))^2]$$

of $f(\mathbf{x}_i)$ and $f(\mathbf{x}_j)$ over all functions f .

Kernels in gaussian processes

- ⊙ The covariance function κ is assumed to be a positive definite (Mercer) kernel.
- ⊙ This means that for any set of distinct points $\mathbf{x}_1, \dots, \mathbf{x}_n$ it must be

$$\sum_{i=1}^n \sum_{j=1}^n c_i c_j \kappa(\mathbf{x}_i, \mathbf{x}_j) > 0$$

for any choice of the constants c_1, \dots, c_n such that not all c_i are equal to 0.

- ⊙ Equivalently, the square **Gram** matrix G defined as

$$G = \begin{pmatrix} \kappa(\mathbf{x}_1, \mathbf{x}_1) & \kappa(\mathbf{x}_1, \mathbf{x}_2) & \cdots & \kappa(\mathbf{x}_1, \mathbf{x}_n) \\ \kappa(\mathbf{x}_2, \mathbf{x}_1) & \kappa(\mathbf{x}_2, \mathbf{x}_2) & \cdots & \kappa(\mathbf{x}_2, \mathbf{x}_n) \\ \cdots & \cdots & \cdots & \cdots \\ \kappa(\mathbf{x}_n, \mathbf{x}_1) & \kappa(\mathbf{x}_n, \mathbf{x}_2) & \cdots & \kappa(\mathbf{x}_n, \mathbf{x}_n) \end{pmatrix}$$

must have positive eigenvalues.

- ⊙ A collection of positive definite kernels is known in the literature and can be constructed by applying suitable rules.

Given a gaussian process $p(f) = \mathcal{GP}(m, \kappa)$, then for any set of items $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$, the distribution of $f(\mathbf{x}_1), \dots, f(\mathbf{x}_n)$ is a gaussian

$$(f(\mathbf{x}_1), \dots, f(\mathbf{x}_n)) \sim \mathcal{N}(\boldsymbol{\mu}(\mathbf{X}) | \Sigma(\mathbf{X}))$$

where

- ⊙ $\boldsymbol{\mu}(\mathbf{X}) = (m(\mathbf{x}_1), \dots, m(\mathbf{x}_n))^T$

- ⊙ $\Sigma(\mathbf{X})$ is the Gram matrix wrt $\mathbf{x}_1, \dots, \mathbf{x}_n$ of a kernel function $\kappa(\mathbf{x}, \mathbf{x}')$

The mean vector - at least initially, with no information from data - is usually assumed to be **0**: different processes are then characterized only by their covariance kernel κ .

- ⊙ For any finite subset $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$ of χ it is possible to sample from $p(f)$ the values of $f(\mathbf{x}_1), \dots, f(\mathbf{x}_m)$ by sampling from $\mathcal{N}(f|\mathbf{0}, \Sigma(\mathbf{X}))$, where, as stated before

$$\mu(\mathbf{X})_i = m(\mathbf{x}_i)$$

$$\Sigma(\mathbf{X})_{ij} = \kappa(\mathbf{x}_i, \mathbf{x}_j)$$


- ⊙ One of the most applied kernel is the RBF kernel

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \sigma^2 e^{-\frac{\|\mathbf{x}_1 - \mathbf{x}_2\|^2}{2\tau^2}}$$

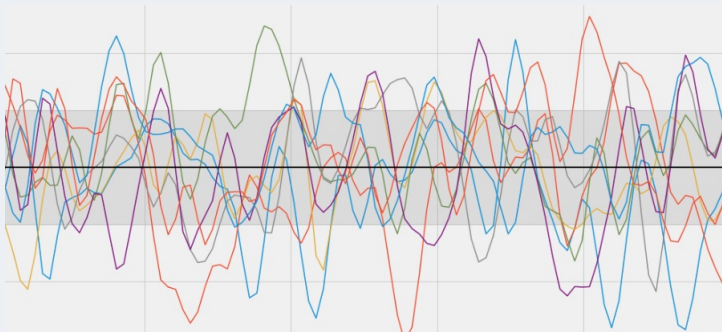
which tends to assign higher covariance between $f(\mathbf{x}_1)$ and $f(\mathbf{x}_2)$ if \mathbf{x}_1 and \mathbf{x}_2 are nearby points.

- ⊙ Functions drawn from a Gaussian process with RBF kernel tend to be smooth (values computed for nearby points tend to have similar values). Smoothing is larger for larger τ .

Larger smoothing




Smaller smoothing



Gaussian process regression: no noise

- ⊙ By the gaussian process definition, \mathbf{f} is distributed as a multivariate gaussian such that the mean of any value \mathbf{x} is $m(\mathbf{x})$ and the covariance of any pair \mathbf{x}, \mathbf{x}' is $\kappa(\mathbf{x}, \mathbf{x}')$
- ⊙ as a consequence, for any finite set of points \mathbf{X} , we have that $\mathbf{f}(\mathbf{X})$ is distributed as a multivariate gaussian with mean $\boldsymbol{\mu}(\mathbf{X})$ defined as $\boldsymbol{\mu}(\mathbf{X})_i = m(\mathbf{x}_i)$ and covariance matrix $\Sigma(\mathbf{X})$, defined as $\Sigma(\mathbf{X})_{i,j} = \kappa(\mathbf{x}_i, \mathbf{x}_j)$



- ⊙ Let us now assume that for a set of points $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T$ the corresponding values $\mathbf{f} = (f(\mathbf{x}_1), \dots, f(\mathbf{x}_n))^T$ are known
- ⊙ that is, we assume that a training set \mathbf{X}, \mathbf{t} is available such that the target values in the training set correspond exactly to the function value $t_i = f(\mathbf{x}_i)$. Note that in the probabilistic model of regression this is not true, since a (gaussian) error is assumed

In this case, for any new set of points \mathbf{X}^* , the joint distribution of $(\mathbf{f}(\mathbf{X}), \mathbf{f}(\mathbf{X}^*))$ is a multivariate gaussian distribution with mean $\boldsymbol{\mu}(\mathbf{X}, \mathbf{X}^*)$ and covariance $\Sigma(\mathbf{X}, \mathbf{X}^*)$


$$\odot \quad \boldsymbol{\mu}(\mathbf{X}, \mathbf{X}^*) = (\boldsymbol{\mu}(\mathbf{X}), \boldsymbol{\mu}(\mathbf{X}^*))^T$$

\odot

$$\Sigma(\mathbf{X}, \mathbf{X}^*) = \begin{pmatrix} \Sigma(\mathbf{X}) & \Sigma(\mathbf{x}, \mathbf{X}) \\ \Sigma(\mathbf{x}, \mathbf{X})^T & \Sigma(\mathbf{X}^*) \end{pmatrix}$$

where $\Sigma(\mathbf{x}, \mathbf{X}) = (\kappa(\mathbf{x}, \mathbf{x}_1), \dots, \kappa(\mathbf{x}, \mathbf{x}_n))$

Gaussian process regression: no noise



The posterior distribution of $\mathbf{f}(\mathbf{X}^*)$, given $\mathbf{X}, \mathbf{X}^*, \mathbf{t}$ can be derived by gaussian distribution properties, and turns out to be a gaussian distribution with mean and covariance defined as


$$\odot \quad \boldsymbol{\mu}^* = \boldsymbol{\mu}(\mathbf{X}^*) + \Sigma(\mathbf{x}, \mathbf{X})\Sigma(\mathbf{X})^{-1}(\mathbf{t} - \boldsymbol{\mu}(\mathbf{X}))$$

$$\odot \quad \Sigma^* = \Sigma(\mathbf{X}^*) - \Sigma(\mathbf{x}, \mathbf{X})\Sigma(\mathbf{X})^{-1}\Sigma(\mathbf{x}, \mathbf{X})^T$$

Gaussian process regression: no noise

In particular, for a single test point \mathbf{x} , we have that the corresponding predictive distribution is

$$m_p(\mathbf{x}|\mathbf{X}, \mathbf{f}) = m(\mathbf{x}) + \Sigma(\mathbf{x}, \mathbf{X})\Sigma(\mathbf{X})^{-1}(\mathbf{t} - \boldsymbol{\mu}(\mathbf{X}))$$
$$\sigma^2 = \kappa_p(\mathbf{x}, \mathbf{x}) = \kappa(\mathbf{x}, \mathbf{x}) - \Sigma(\mathbf{x}, \mathbf{X})\Sigma(\mathbf{X})^{-1}\Sigma(\mathbf{x}, \mathbf{X})^T$$



In this case, an **interpolation** of the given values is performed: $f(\mathbf{x}_i) = t_i$ for all possible functions, sampled from $f(\mathbf{x}|\mathbf{X}, \mathbf{f})$.

In fact, for all $\mathbf{x}_i \in \mathbf{X}$,

- ⊙ $\Sigma(\mathbf{x}_i, \mathbf{X})$ is the i -th row of $\Sigma(\mathbf{X})$: since $\Sigma(\mathbf{X})\Sigma(\mathbf{X})^{-1} = \mathbf{I}$ by definition, the product $\Sigma(\mathbf{x}_i, \mathbf{X})\Sigma(\mathbf{X})^{-1}$ must be equal to the i -th row of \mathbf{I} , that is an array C_i such that $C_i[i] = 1$ and $C_i[j] = 0$ otherwise
- ⊙ as a consequence,

$$\begin{aligned}m(\mathbf{x}_i|\mathbf{X}, \mathbf{f}) &= m(\mathbf{x}_i) + (t_i - m(\mathbf{x}_i)) = t_i \\ \sigma^2 &= \kappa(\mathbf{x}_i, \mathbf{x}_i) - \kappa(\mathbf{x}_i, \mathbf{x}_i) = 0\end{aligned}$$

Let us now assume, as usual, that $p(t_i|f, \mathbf{x}_i) = \mathcal{N}(f(\mathbf{x}_i), \sigma_f^2)$


That is, the value t_i observed for variable \mathbf{x}_i differs from the one obtained as $f(\mathbf{x}_i)$ by a gaussian and independent noise

$$t_i = f(\mathbf{x}_i) + \varepsilon \qquad p(\varepsilon) = \mathcal{N}(\varepsilon|0, \sigma_f^2)$$

that is, $p(\mathbf{t}|\mathbf{f}) = \mathcal{N}(\mathbf{t}|\mathbf{f}, \sigma^2\mathbf{I})$

Gaussian process regression: gaussian noise

- ⊙ \mathbf{f} is now distributed as a multivariate gaussian with known mean $\boldsymbol{\mu}(\mathbf{X}) = (m(\mathbf{x}_1), \dots, m(\mathbf{x}_n))^T$ and covariance matrix $\hat{\Sigma}(\mathbf{X}) = \Sigma(\mathbf{X}) + \sigma_f^2 \mathbf{I}$, defined by κ and σ_f^2



- ⊙ Let us now assume that a training set \mathbf{X}, \mathbf{t} is available such that the target values in the training set correspond approximately to the function value $t_i = f(\mathbf{x}_i) + \varepsilon$.
- ⊙ In this case, for any new set of points \mathbf{X}^* , the joint distribution of $(\mathbf{t}, \mathbf{f}(\mathbf{X}^*))$ is a multivariate gaussian distribution with mean $\boldsymbol{\mu}(\mathbf{X}, \mathbf{X}^*)$ and covariance $\Sigma(\mathbf{X}, \mathbf{X}^*)$
- ⊙ $\boldsymbol{\mu}(\mathbf{X}, \mathbf{X}^*) = (\boldsymbol{\mu}(\mathbf{X}), \boldsymbol{\mu}(\mathbf{X}^*))^T$
- ⊙


$$\Sigma(\mathbf{X}, \mathbf{X}^*) = \begin{pmatrix} \hat{\Sigma}(\mathbf{X}) & \Sigma(\mathbf{x}, \mathbf{X}) \\ \Sigma(\mathbf{x}, \mathbf{X})^T & \Sigma(\mathbf{X}^*) \end{pmatrix}$$

Gaussian process regression: gaussian noise

The posterior distribution of $\mathbf{f}(\mathbf{X}^*)$, given $\mathbf{X}, \mathbf{X}^*, \mathbf{t}$ can be again derived by the gaussian distribution properties, and turns out again to be a gaussian distribution with mean and covariance defined as

$$\odot \quad \boldsymbol{\mu}^* = \boldsymbol{\mu}(\mathbf{X}^*) + \Sigma(\mathbf{x}, \mathbf{X}) \hat{\Sigma}(\mathbf{X})^{-1} (\mathbf{t} - \boldsymbol{\mu}(\mathbf{X}))$$


$$\odot \quad \Sigma^* = \Sigma(\mathbf{X}^*) - \Sigma(\mathbf{x}, \mathbf{X}) \hat{\Sigma}(\mathbf{X})^{-1} \Sigma(\mathbf{x}, \mathbf{X})^T$$



Gaussian process regression: gaussian noise

In particular, for a single test point \mathbf{x} , we have now that the corresponding predictive distribution is

$$m_p(\mathbf{x}|\mathbf{X}, \mathbf{f}) = m(\mathbf{x}) + \Sigma(\mathbf{x}, \mathbf{X})\hat{\Sigma}(\mathbf{X})^{-1}(\mathbf{t} - \boldsymbol{\mu}(\mathbf{X}))$$
$$\sigma^2 = \kappa_p(\mathbf{x}, \mathbf{x}) = \kappa(\mathbf{x}, \mathbf{x}) - \Sigma(\mathbf{x}, \mathbf{X})\hat{\Sigma}(\mathbf{X})^{-1}\Sigma(\mathbf{x}, \mathbf{X})^T$$



Estimating kernel parameters


The predictive performance of gaussian processes depends exclusively on the suitability of the chosen kernel.

Let us consider the case of an RBF kernel. Then,

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \sigma_f^2 e^{-\frac{1}{2}(\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{M}(\mathbf{x}_i - \mathbf{x}_j)} + \sigma_y^2 \delta_{ij}$$

\mathbf{M} can be defined in several ways: the simplest one is $\mathbf{M} = l^{-2} \mathbf{I}$.

Even in this simple case, varying the values of σ_f, σ_y, l returns quite different results.



(figure from K.Murphy “Machine learning: a probabilistic perspective” p. 519, with (l, σ_f, σ_y) equal to $(1, 1, 0.1)$, $(0.3, 1.08, 0.00005)$, $(3.0, 1.16, 0.89)$)

Kernel parameters can be estimated, as usual, through grid search and (cross-)validation.

A different, more efficient approach relies on maximizing the marginal likelihood

$$p(\mathbf{t}|\mathbf{X}) = \int p(\mathbf{t}|\mathbf{f}, \mathbf{X})p(\mathbf{f}|\mathbf{X})d\mathbf{f} = \int \mathcal{N}(\mathbf{f}|\mathbf{0}, \Sigma(\mathbf{X})) \prod_{i=1}^m \mathcal{N}(t_i|f_i, \sigma_t^2) d\mathbf{f}$$